

Advanced Spring MVC for High-Traffic Applications

Q1. What are some advanced features or techniques in Spring MVC that are useful for high-traffic applications?

A: Two useful techniques are **asynchronous request processing** (using `DeferredResult` or `Callable`) to prevent thread exhaustion during slow I/O. Another technique is **Content Negotiation**, which automatically handles the correct serialization (JSON/XML) based on client request headers, keeping the code clean and efficient.

Q2. How can caching be implemented in Spring MVC?

A: Caching is implemented using the **@Cacheable** annotation on service methods, which tells Spring to store the result of that method call in a cache (like Redis). For HTTP caching, you use the **ETag** and **Last-Modified** headers in your controller to instruct the client's browser or proxy server to use a locally cached version of the response.

Q3. What are the strategies for asynchronous processing in Spring MVC?

A: The primary strategies are returning **DeferredResult** or **Callable** objects from the controller. This allows the primary thread (the servlet thread) to be released immediately back to the pool while the task runs in a separate worker thread. When the task finishes, the result is delivered back to the client via the original request context.

Q4. How can you scale a Spring MVC application horizontally?

A: You scale horizontally by running multiple identical instances of the application behind a **load balancer**. To make this work, the application must be designed to be **stateless**, meaning no user data (session data) is stored locally on the server. All shared data must be managed externally in a central, distributed store like Redis.

Q5. Explain how you can handle static resources at scale (CDN, cache headers, versioning).

A: To handle static resources at scale, you place them on a **Content Delivery Network (CDN)** to serve files closer to users. You ensure resources have long **Cache-Control** headers so browsers cache them for extended periods. You use **versioning** (e.g., appending a hash to the filename) so browsers pull a new file only when the content actually changes.

Q6. How do you tune the DispatcherServlet for high throughput?

A: You tune the DispatcherServlet by configuring the **embedded servlet container's thread pool** size to match the available CPU cores and expected load. You should also ensure that the application uses non-blocking resources (like non-blocking I/O) as much as possible, preventing threads from being stuck waiting.

Q7. How does Spring MVC behave differently in a microservices architecture vs monolith?

A: In a **monolith**, Spring MVC handles both the web layer and the view rendering. In a **microservices** architecture, Spring MVC is typically used only for the **REST API** layer. It acts as a controller, handling requests and calling business services, but it usually returns only data (JSON), delegating view rendering to a separate frontend.

Q8. How do you integrate Spring MVC with enterprise systems (Kafka, Redis, external auth providers)?

A: Integration is achieved using **Dependency Injection (DI)**. Controllers call services, and services use Spring clients: **Spring Data Redis** for caching, **Spring Kafka** for messaging (queues), and **Spring Security** (configured for OAuth2/JWT) for external authentication. All these clients are injected as beans.

Q9. How to avoid performance bottlenecks in view rendering and serialization?

A: Avoid bottlenecks by using efficient view templates (like **Thymeleaf**) and ensuring templates do not execute complex database queries. For **serialization** (JSON), ensure you use efficient DTOs (Data Transfer Objects) that contain only the necessary fields, minimizing the data transferred over the network.

Q10. What is an optimal approach for handling file uploads at scale? (streaming upload, chunking)

A: The optimal approach is **streaming upload**. Instead of buffering the entire file into the server's heap memory, you stream the data directly to external persistent storage (like S3) using the `InputStream`. For very large files, **chunking** is used, where the client breaks the file into small, manageable pieces that are uploaded sequentially and reassembled on the server side.

Q11. How do you integrate async request processing with `WebAsyncTask` / `DeferredResult`?

A: You integrate by having the controller method return a **DeferredResult** object. When the request hits the controller, the servlet thread is released immediately. The `DeferredResult` is handed to a worker thread which holds it until the long task finishes. Once complete, the worker thread sets the result onto the `DeferredResult`, which sends the response back to the client.

Q12. How does Spring MVC compare to Spring WebFlux in high concurrency scenarios?

A: **Spring MVC** (blocking model) requires a larger thread pool because each request uses a thread for its full duration (thread-per-request). This limits concurrency. **Spring WebFlux** (non-blocking model) can handle much higher concurrency with a smaller number of threads because threads are not held while waiting for I/O operations to complete.

Q13. What architectural patterns are needed to prevent thread exhaustion in MVC apps?

A: The main pattern is the **Asynchronous Request Processing** pattern (using `DeferredResult` or `Callable`) to prevent threads from being tied up waiting for slow I/O. Other patterns include using a **Circuit Breaker** (Resilience4J) on external dependencies, which immediately returns a fallback instead of waiting for a failing service.

Q14. In high-traffic apps, how do you reduce session load and dependency? (stateless design, tokenization)

A: You reduce session load by transitioning to a **stateless design**, where no user-specific data is stored on the server side. Authentication is handled using **tokenization** (JWTs), which are self-contained and sent with every request. If session data is absolutely necessary, it should be externalized to a highly scalable store like **Redis**.